

Linear Machine Learning

Janis Keuper

Basic Types of Machine Learning Algorithms

Supervised Learning

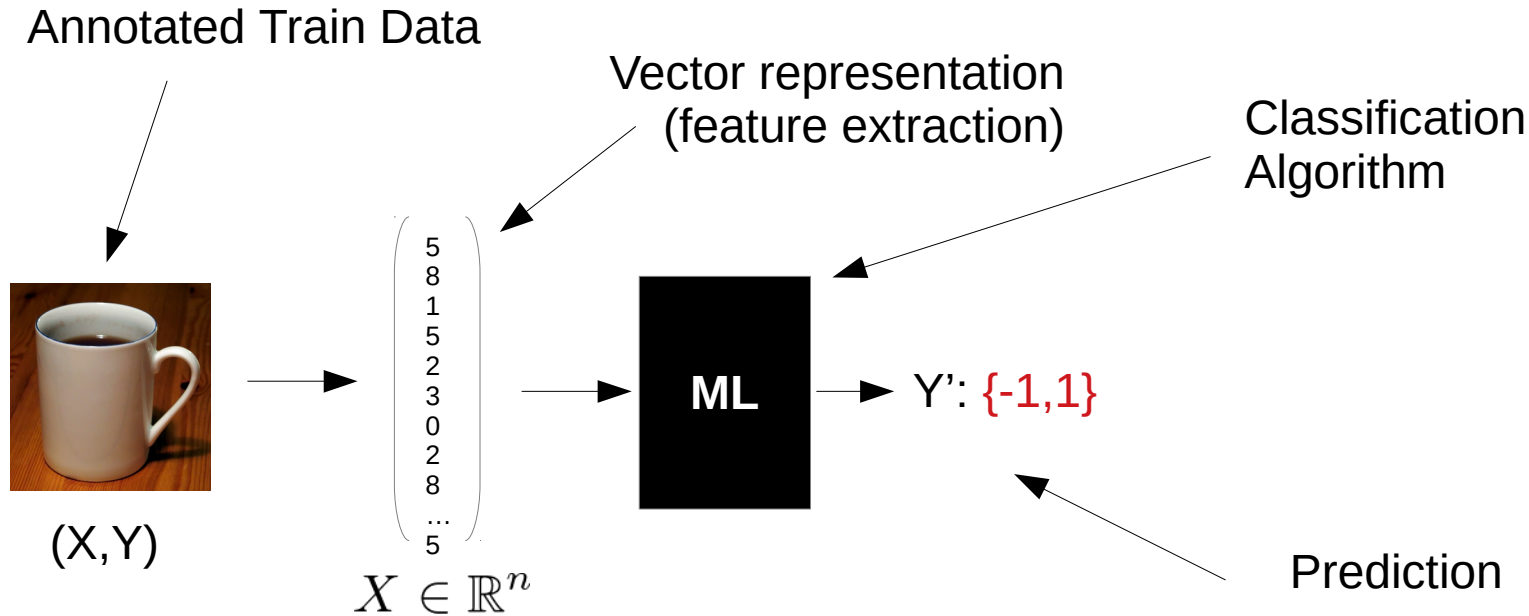
Unsupervised Learning

Reinforcement Learning

- Labeled data
- Direct and quantitative evaluation
- Learn model from „ground truth“ examples
- Predict unseen examples

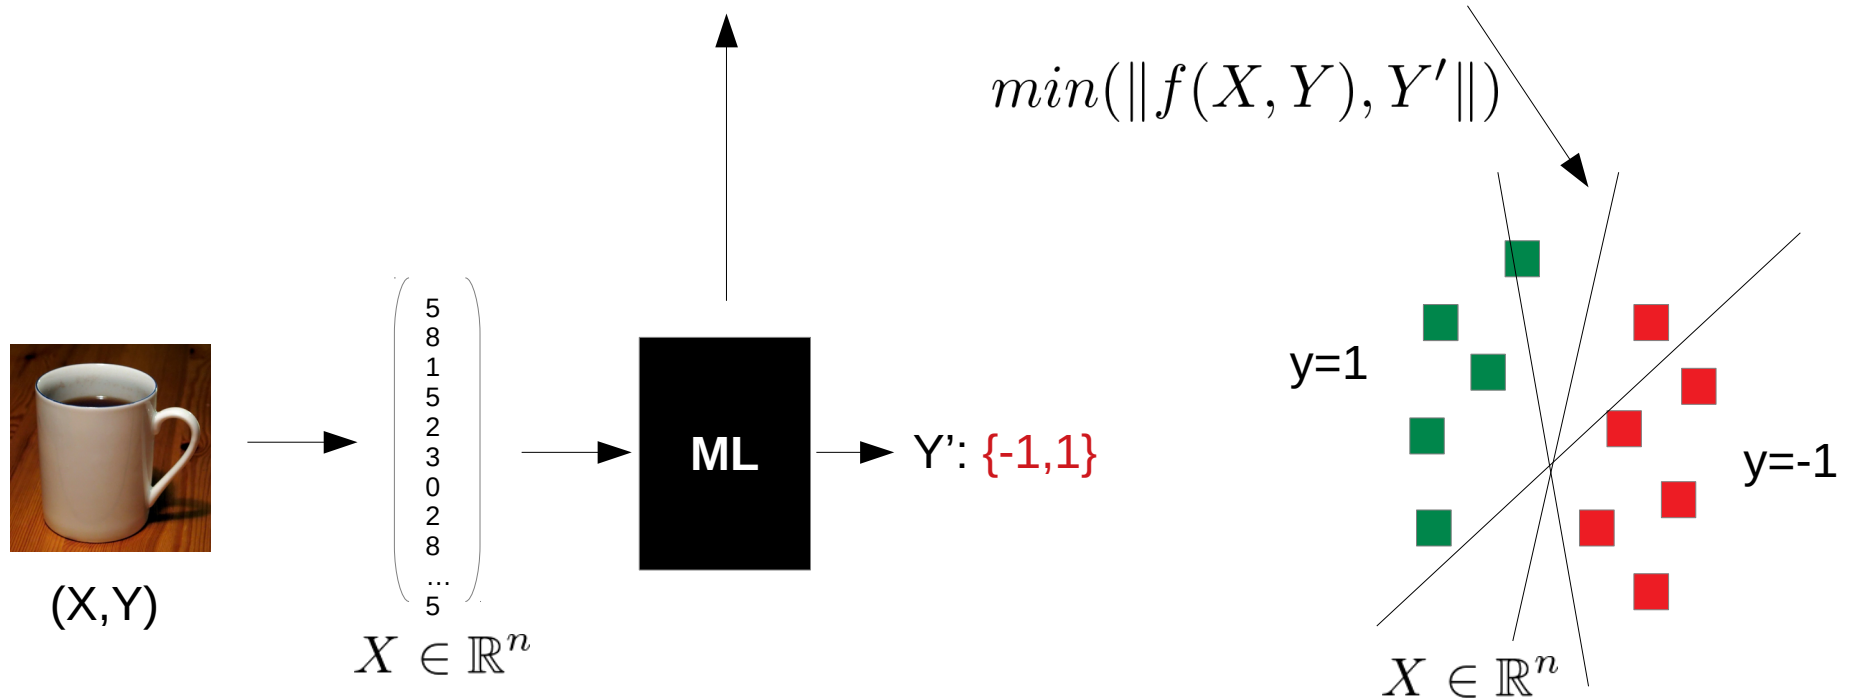
Recall: Classification

Supervised Learning: Annotated Training Data



Recall: Classification

LEARNING: is a optimization problem → Finding the best function separating



Recall: Linear Classifier

A Simple Linear Model: **binary** classification


Parameterization of prediction function f
with d -dimensional data as:

$$f(x) = y' = w^T x = \sum_{j=0}^d w_j x_j$$

With data samples $x \in \mathbb{R}^d$

Model parameters $w \in \mathbb{R}^d$

Model: hyper plane




Limitations of Linear Models

(Obviously), linear models have limitations

Consider this very simple binary classification
Example:

How to separate “green” from “red”
with a linear Model (= hyper plane)?

Simple counter example



$$x \in \mathbb{R}^d$$

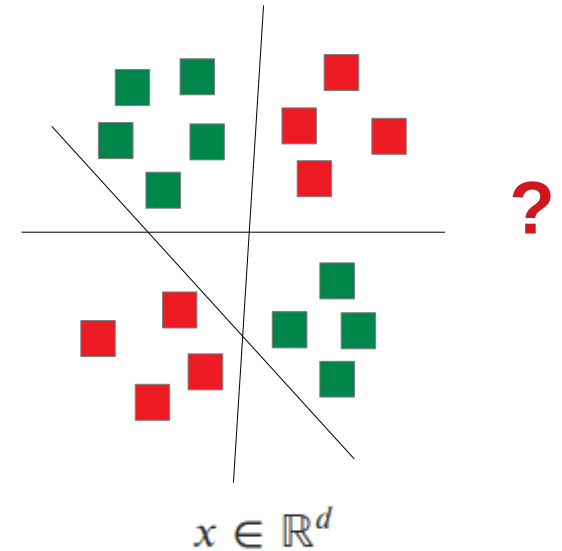
Limitations of Linear Models

(Obviously), linear models have limitations

Consider this very simple binary classification
Example:

How to separate “green” from “red”
with a linear Model (= hyper plane)?

Simple counter example



Limitations of Linear Models

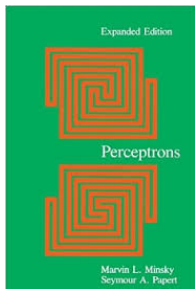
(Obviously), linear models have limitations

Consider this very simple binary classification
Example:

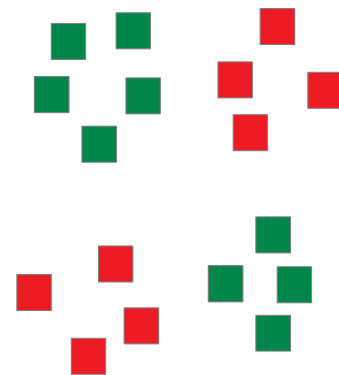
Simple counter example

→ known as “X-Or” Problem

→ one reason for the so-called “AI Winter”



Caused by the Minsky book
On the shortcomings of the
First neural networks...

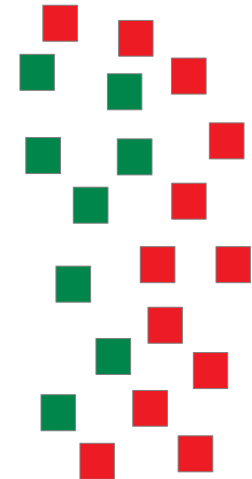
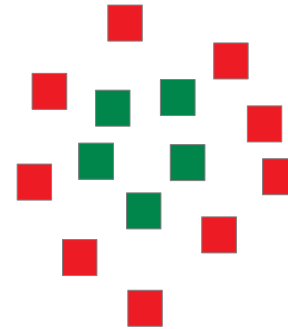
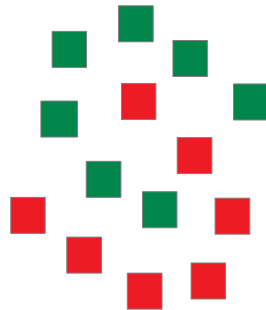
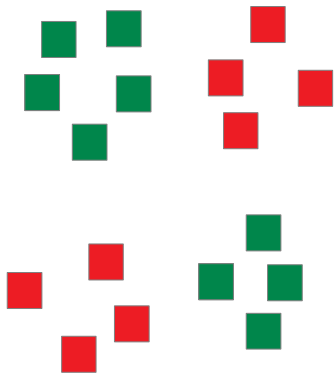


$x \in \mathbb{R}^d$

Limitations of Linear Models

(Obviously), linear models have limitations

More simple (binary 2D) examples:



Why are linear model working at all?

Why are linear model working at all?

- **very high dimensional feature spaces often time allow linear models to Separate the data**
- **very simple (linear) model even can be of advantage in theses settings:**
 - **“curse of dimensionality”**
 - **Avoid overfitting**

Adding non-linearity

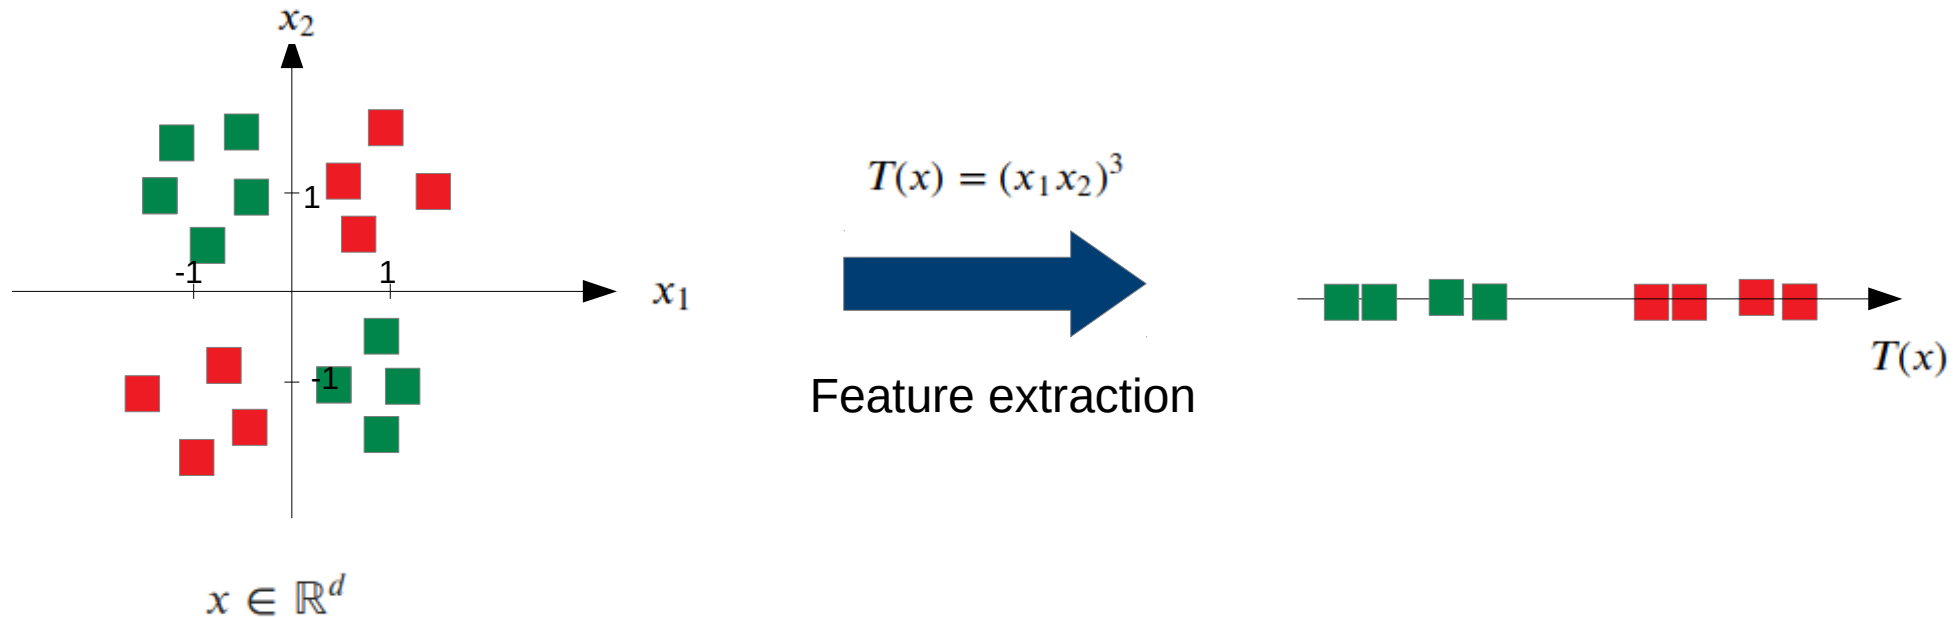


Three canonical ways (we already saw two of them):

Adding non-linearity

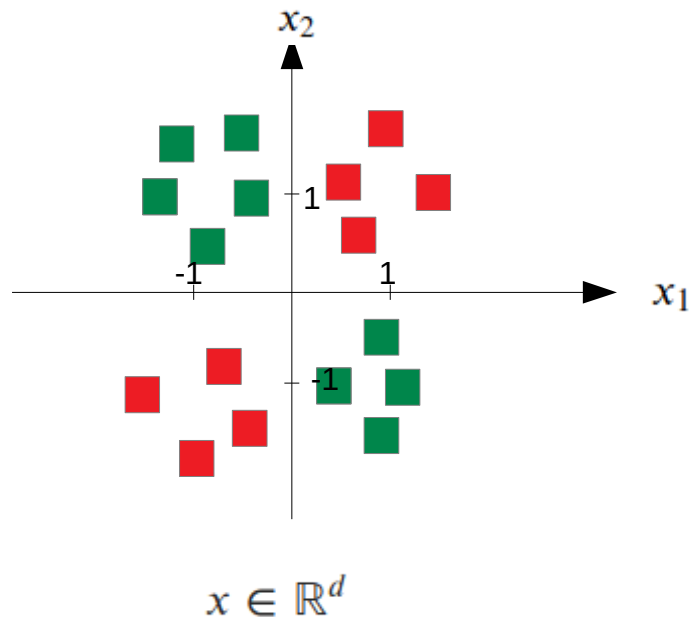
Three canonical ways:

1. extracting non-linear features:




Three canonical ways:


1. extracting non-linear features:



$T(x) = (x_1 x_2)^3$



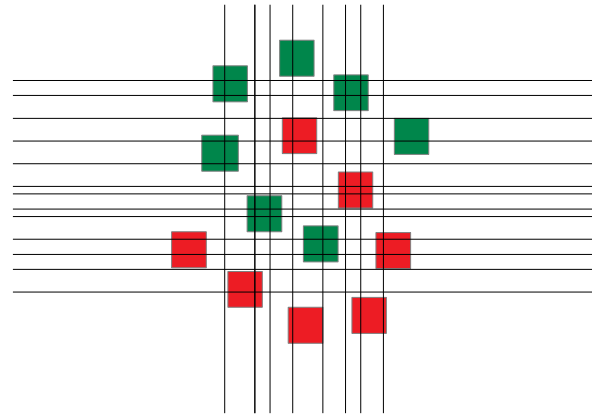
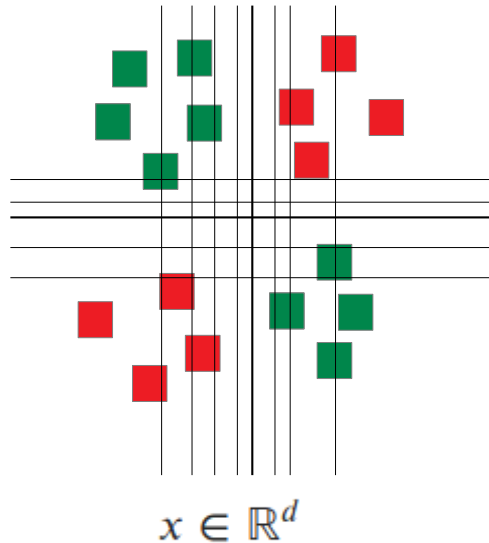
Feature extraction



NOTE: just an example, usually not
That simple for real data.

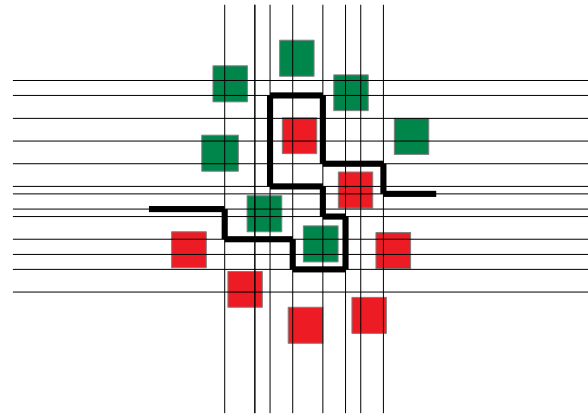
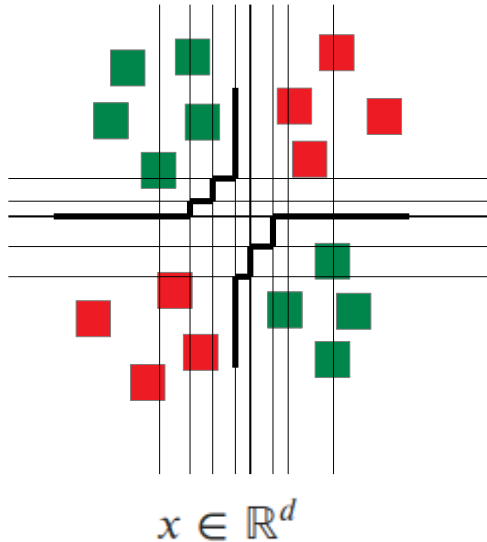
Three canonical ways:

2. use ensembles of linear models (like Random Forrest)




Three canonical ways:

2. use ensembles of linear models → approximation of non-linear models by piece-wise linear models



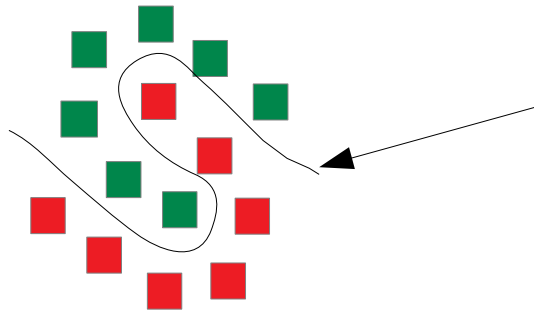
Three canonical ways:

3. use non-linear functions



Three canonical ways:


3. use non-linear functions




How to parameterize the non-linear model?

Adding non-linearity to our simple linear classifier

$$f(x) = y' = w^T x = \sum_{j=0}^d w_j x_j \quad \xrightarrow{\text{Train by solving optimization}} \quad \arg \min_w \sum_{i=0}^N L(y_i, w^T x_i)$$



$$f(x) = \phi(w^T x)$$



Step I: add a very simple element-wise non-linear mapping.

(like in the previous feature extraction example)

Adding non-linearity to our simple linear classifier

$$f(x) = y' = w^T x = \sum_{j=0}^d w_j x_j$$



$$f(x) = \phi(w^T x)$$



What are good choices for these functions?

Adding non-linearity to our simple linear classifier

$$f(x) = y' = w^T x = \sum_{j=0}^d w_j x_j$$




$$f(x) = \phi(w^T x)$$



What are good choices for these functions?

Properties:

- Between 0 and 1 → pseudo probability interpretation
- Stable range of output → gradient optimization



Common choices:

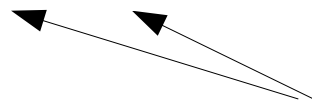
- **Sigmoid function**
- Tanh
- ...

Adding non-linearity to our simple linear classifier

$$f(x) = y' = w^T x = \sum_{j=0}^d w_j x_j$$



$$f(x) = \phi_3(w_3 \phi_2(W_2 \phi_1(W_1 x)))$$



W are now Matrices to produce vector outputs

Step II: concatenate several of these operations

(like we do in the ensemble approach)


Adding non-linearity

Let's display this in a slightly different way (no change in math formulation!)

$$f(x) = y' = w^T x = \sum_{j=0}^d w_j x_j$$



$$f(x) = \phi_3(w_3 \phi_2(W_2 \phi_1(W_1 x)))$$



Matrix/Vector Mult

Element wise non-linear


Adding non-linearity

Let's display this in a slightly different way (no change in math formulation!)

$$f(x) = y' = w^T x = \sum_{j=0}^d w_j x_j$$



$$f(x) = \phi_3(w_3 \phi_2(W_2 \phi_1(W_1 x)))$$



Note: there is a theoretical prove that we need only two concatenations to approximate any smooth function if the W are large enough!

Matrix/Vector Mult

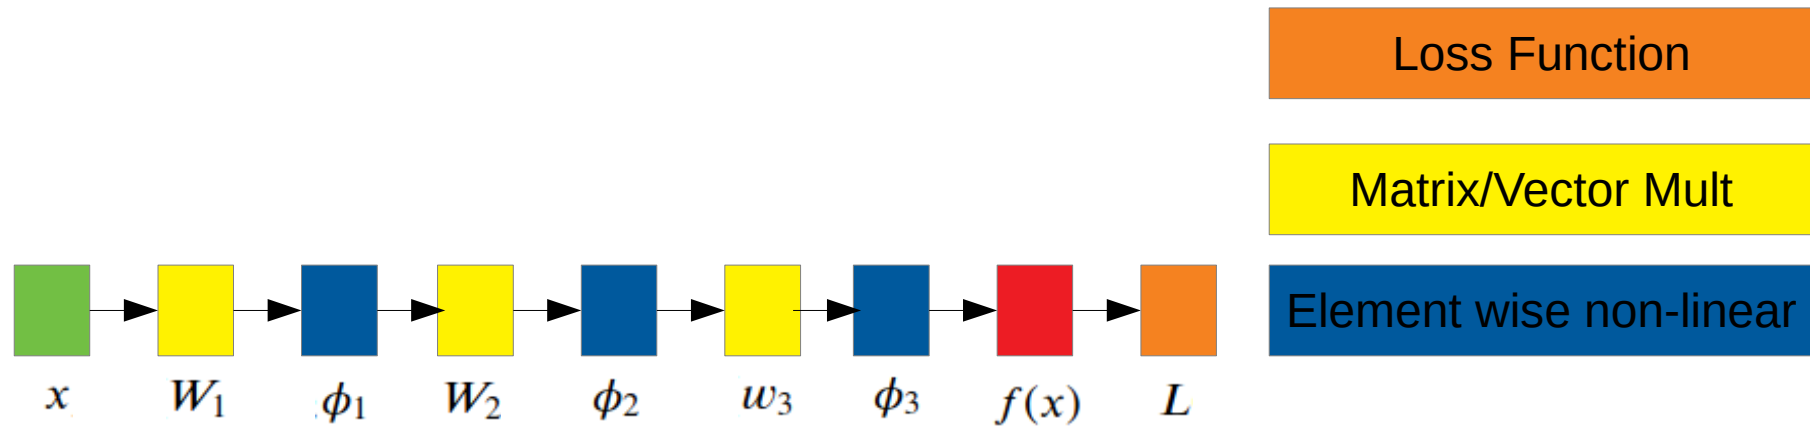
Element wise non-linear

Adding non-linearity

For training (optimization), we need to add loss function

→ same approach as in the linear case:

$$\arg \min_w \sum_{i=0}^N L(y_i, f(x))$$




Adding non-linearity

For training (optimization), we need to add loss function

→ same approach as in the linear case:

$$\arg \min_w \sum_{i=0}^N L(y_i, f(x))$$

→ solve by gradient descent optimization



Loss Function

Matrix/Vector Mult

Element wise non-linear

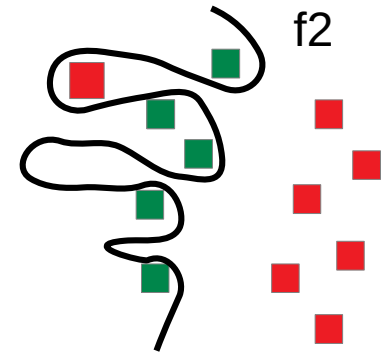
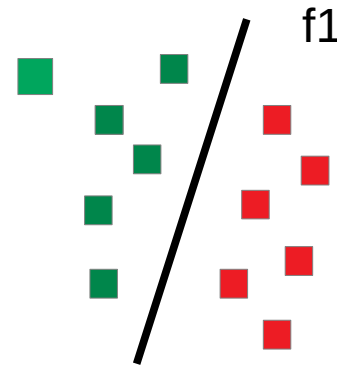
Adding non-linearity

Recall **OVERFITTING**

Model “to close” to train data

With non-linear model much more likely to happen in practice.

→ we need to work against this...



Adding regularization term to the Loss function

→ here L2 regularization:

$$\arg \min_w \sum_{i=0}^N L(y_i, f(x)) + \lambda \sum_j w_j^2$$

All parameters to be learned

Scalar hyper parameter: impact of regularization

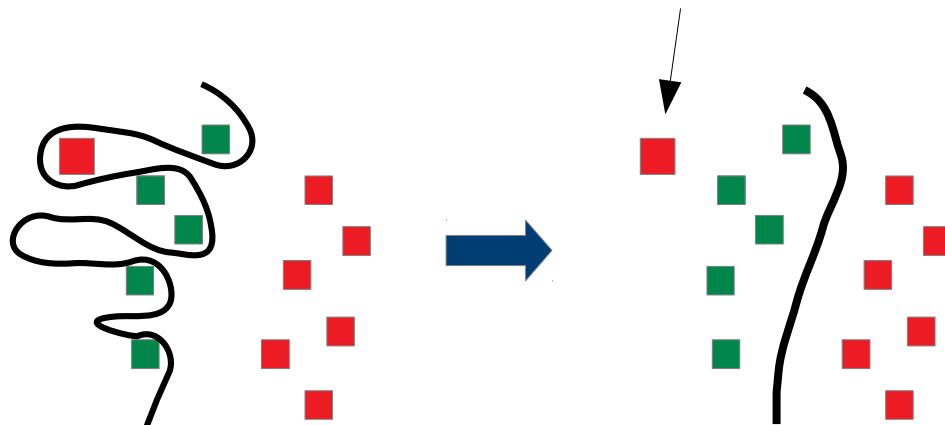
Adding regularization term to the Loss function

→ here L2 regularization:

$$\arg \min_w \sum_{i=0}^N L(y_i, f(x)) + \lambda \sum_j w_j^2$$

→ L1 regularization:

$$\arg \min_w \sum_{i=0}^N L(y_i, f(x)) + \lambda \sum_j |w_j|$$



Regularization will punish high parameter values

→ smoother model

→ training errors allowed !

Support Vector Machines



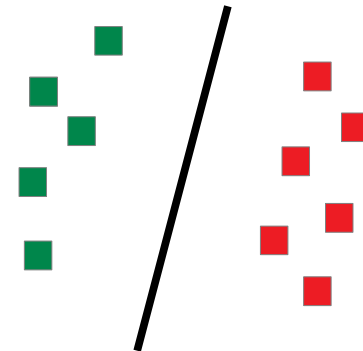
INSTITUTE FOR MACHINE
LEARNING AND ANALYTICS

Support Vector Machines

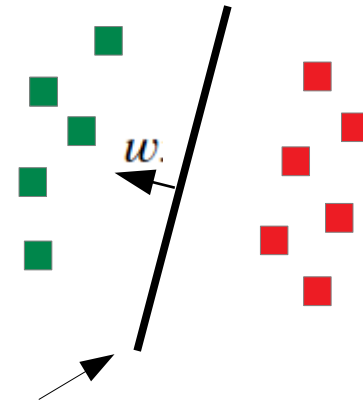


- Invented in the mid 90s by Vapnik
- Classification and Regression
- State of the Art ML Algorithm of the pre Deep Learning era

- Invented in the mid 90s by Vapnik
- Classification and Regression
- State of the Art ML Algorithm of the pre Deep Learning era
- **Basic model:**
 - Support only two classes $\{-1,1\}$
 - Linear classification




- Invented in the mid 90s by Vapnik
- Classification and Regression
- State of the Art ML Algorithm of the pre Deep Learning era
- **Basic model:**
 - Support only two classes $\{-1,1\}$
 - Linear classification



Parameterization: $wx - b = 0$

What is the difference compared to previous formulations?

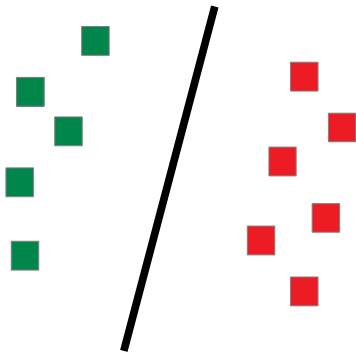


Standard linear model:

- loss only on accuracy
- many solutions

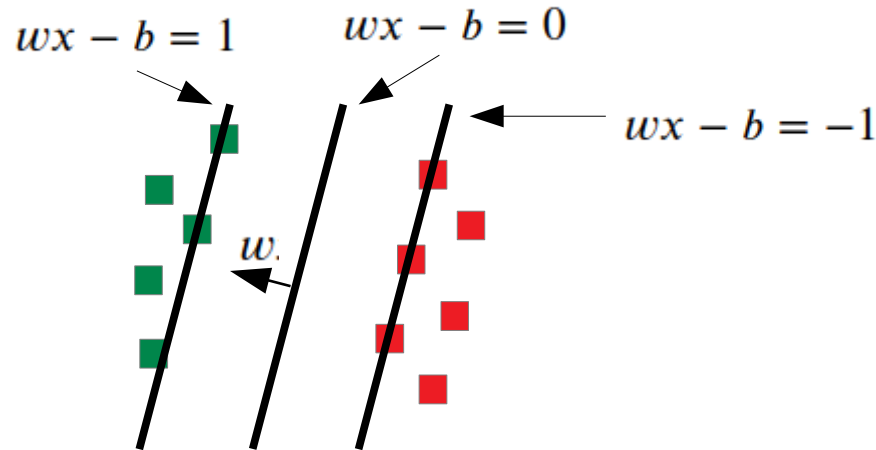
$$\arg \min_w \sum_{i=0}^N L(y_i, w^T x_i)$$

What is the difference compared to previous formulations?



Standard linear model:

- loss only on accuracy
- many solutions



New optimization problem

→ “Max Margin”: $\frac{2}{\|w\|}$

→ only one solution, convex optimization problem

What is the difference compared to previous formulations?

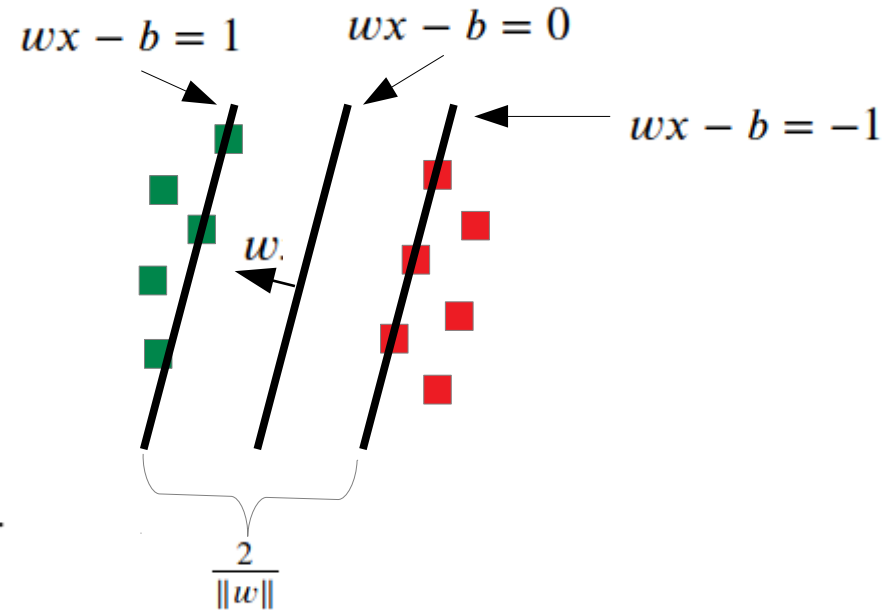
New optimization problem

- maximize “Margin”:
- equals minimizing the uncertainty

$$\arg \min_w \sum_{i=0}^N \xi_i + \lambda \|w\|^2$$

subject to $y_i(w \cdot x_i - b) \geq 1 - \xi_i$ and $\xi_i \geq 0$, for all i .

$$\zeta_i = \max(0, 1 - y_i(w \cdot x_i - b))$$



What is the difference compared to previous formulations?

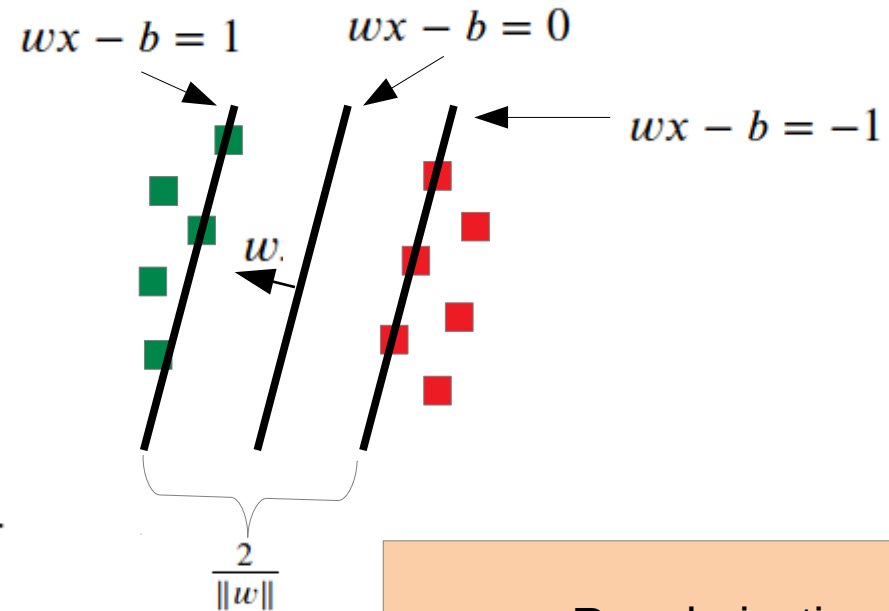
New optimization problem

- maximize “Margin”:
- equals minimizing the uncertainty

$$\arg \min_w \sum_{i=0}^N \xi_i + \lambda \|w_i\|^2$$

subject to $y_i(w \cdot x_i - b) \geq 1 - \xi_i$ and $\xi_i \geq 0$, for all i .

$$\zeta_i = \max(0, 1 - y_i(w \cdot x_i - b))$$



Regularization

What is the difference compared to previous formulations?


New optimization problem

- maximize “Margin”:
- equals minimizing the uncertainty

$$\arg \min_w \sum_{i=0}^N \xi_i + \lambda \|w_i\|^2$$

subject to $y_i(w \cdot x_i - b) \geq 1 - \zeta_i$ and $\zeta_i \geq 0$, for all i .

$$\zeta_i = \max(0, 1 - y_i(w \cdot x_i - b))$$



Max Margin conditions

Loss based on conditions

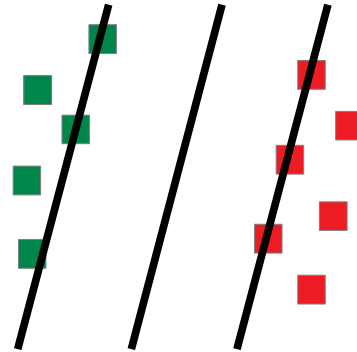
Dual Formulation of the optimization Problem

Via Lagrange dual function, leads to quadratic optimization problem (convex)

$$\vec{w} = \sum_{i=1}^n c_i y_i \vec{x}_i$$

$$\text{maximize } f(c_1 \dots c_n) = \sum_{i=1}^n c_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n y_i c_i (x_i \cdot x_j) y_j c_j,$$

$$\text{subject to } \sum_{i=1}^n c_i y_i = 0, \text{ and } 0 \leq c_i \leq \frac{1}{2n\lambda} \text{ for all } i.$$




Dual Formulation of the optimization Problem

Via Lagrange dual function, leads to quadratic optimization problem (convex)

$$\vec{w} = \sum_{i=1}^n c_i y_i \vec{x}_i$$

$$\text{maximize } f(c_1 \dots c_n) = \sum_{i=1}^n c_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n y_i c_i (x_i \cdot x_j) y_j c_j,$$

$$\text{subject to } \sum_{i=1}^n c_i y_i = 0, \text{ and } 0 \leq c_i \leq \frac{1}{2n\lambda} \text{ for all } i.$$



Re-write model as linear combination
of all weighted (c) data points

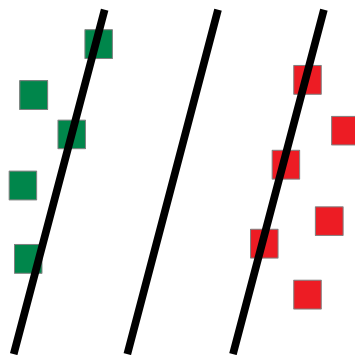
Dual Formulation of the optimization Problem

Via Lagrange dual function, leads to quadratic optimization problem (convex)

$$\vec{w} = \sum_{i=1}^n c_i y_i \vec{x}_i$$

$$\text{maximize } f(c_1 \dots c_n) = \sum_{i=1}^n c_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n y_i c_i (x_i \cdot x_j) y_j c_j,$$

$$\text{subject to } \sum_{i=1}^n c_i y_i = 0, \text{ and } 0 \leq c_i \leq \frac{1}{2n\lambda} \text{ for all } i.$$



Regularization


Re-write model as linear combination
of all weighted (c) data points

How many Data Points do we need to define a (hyper) plane?

$$\vec{w} = \sum_{i=1}^n c_i y_i \vec{x}_i$$

$$\text{maximize } f(c_1 \dots c_n) = \sum_{i=1}^n c_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n y_i c_i (x_i \cdot x_j) y_j c_j,$$

$$\text{subject to } \sum_{i=1}^n c_i y_i = 0, \text{ and } 0 \leq c_i \leq \frac{1}{2n\lambda} \text{ for all } i.$$

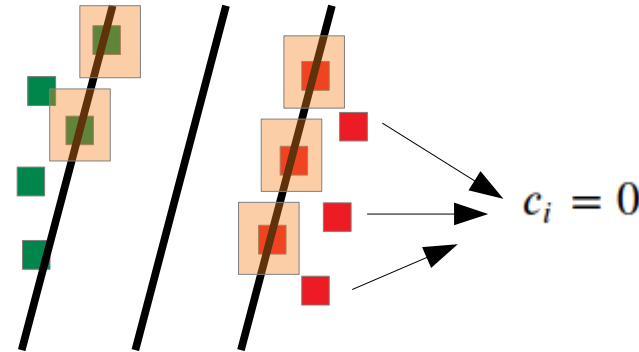


How many Data Points do we need to define a (hyper) plane?

$$\vec{w} = \sum_{i=1}^n c_i y_i \vec{x}_i$$

$$\text{maximize } f(c_1 \dots c_n) = \sum_{i=1}^n c_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n y_i c_i (x_i \cdot x_j) y_j c_j,$$

$$\text{subject to } \sum_{i=1}^n c_i y_i = 0, \text{ and } 0 \leq c_i \leq \frac{1}{2n\lambda} \text{ for all } i.$$




Most point will not contribute.
Only “**Support Vectors**” will.

SVM Model: All Support Vectors

$$\vec{w} = \sum_{i=1}^n c_i y_i \vec{x}_i$$

$$\text{maximize } f(c_1 \dots c_n) = \sum_{i=1}^n c_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n y_i c_i (x_i \cdot x_j) y_j c_j,$$

$$\text{subject to } \sum_{i=1}^n c_i y_i = 0, \text{ and } 0 \leq c_i \leq \frac{1}{2n\lambda} \text{ for all } i.$$

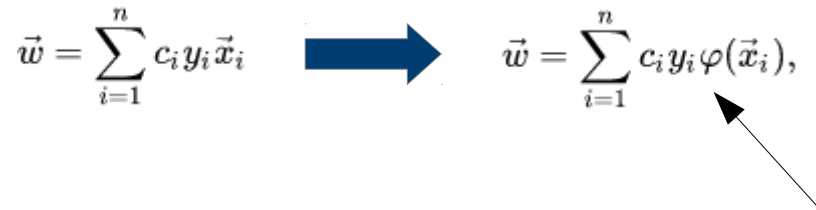


Most point will not contribute.
Only “**Support Vectors**” will.

Non Linear SVMs:

→ follow same strategy as before and add simple non-linear function

At the formulation of the model normal:

$$\vec{w} = \sum_{i=1}^n c_i y_i \vec{x}_i \quad \longrightarrow \quad \vec{w} = \sum_{i=1}^n c_i y_i \varphi(\vec{x}_i),$$
A large blue arrow points from the first equation to the second. A thin black arrow points from the phi symbol in the second equation to the right.

Non Linear SVMs:

→ follow same strategy as before and add simple non-linear function

At the formulation of the model normal:

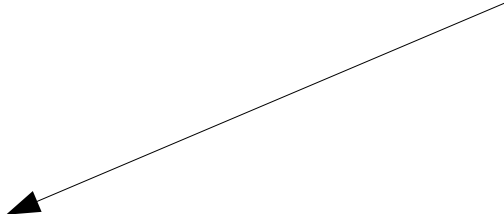
$$\vec{w} = \sum_{i=1}^n c_i y_i \vec{x}_i \quad \longrightarrow \quad \vec{w} = \sum_{i=1}^n c_i y_i \varphi(\vec{x}_i),$$



Insert into dual formulation

$$\text{maximize } f(c_1 \dots c_n) = \sum_{i=1}^n c_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n y_i c_i (\varphi(\vec{x}_i) \cdot \varphi(\vec{x}_j)) y_j c_j$$

“Kernel Trick”: replace explicit non-linear function by *kernel*

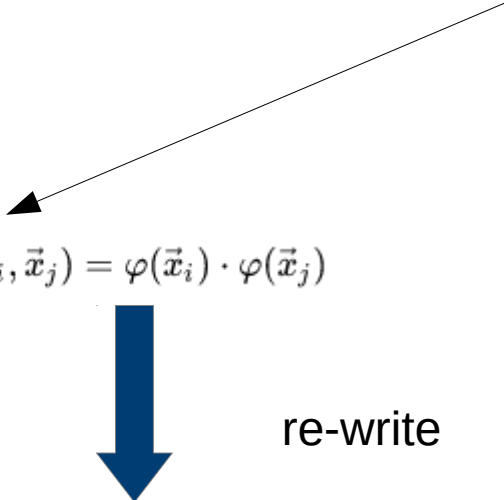

$$k(\vec{x}_i, \vec{x}_j) = \varphi(\vec{x}_i) \cdot \varphi(\vec{x}_j)$$



Always a dot-product in some space

$$\text{maximize } f(c_1 \dots c_n) = \sum_{i=1}^n c_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n y_i c_i (\varphi(\vec{x}_i) \cdot \varphi(\vec{x}_j)) y_j c_j$$

“Kernel Trick”: replace explicit non-linear function by **kernel**


$$k(\vec{x}_i, \vec{x}_j) = \varphi(\vec{x}_i) \cdot \varphi(\vec{x}_j)$$

re-write

$$\begin{aligned} \text{maximize } f(c_1 \dots c_n) &= \sum_{i=1}^n c_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n y_i c_i (\varphi(\vec{x}_i) \cdot \varphi(\vec{x}_j)) y_j c_j \\ &= \sum_{i=1}^n c_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n y_i c_i k(\vec{x}_i, \vec{x}_j) y_j c_j \end{aligned}$$

Popular Kernels:

Polynomial (of degree d)

$$k(\vec{x}_i, \vec{x}_j) = (\vec{x}_i \cdot \vec{x}_j)^d$$

Gauss (or RBF)

$$k(\vec{x}_i, \vec{x}_j) = \exp(-\gamma \|\vec{x}_i - \vec{x}_j\|^2)$$

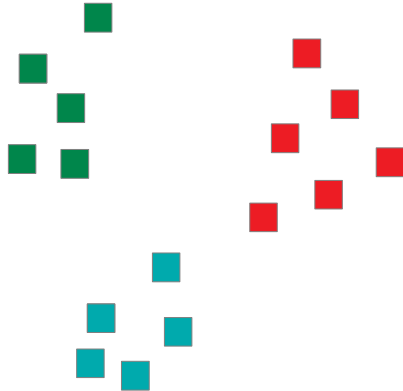
SVM Inference:

→ evaluate kernel with all Support Vectors and take the sign


$$\vec{z} \mapsto \text{sgn}(\vec{w} \cdot \varphi(\vec{z}) - b) = \text{sgn}\left(\left[\sum_{i=1}^n c_i y_i k(\vec{x}_i, \vec{z})\right] - b\right).$$

$$\begin{aligned} b = \vec{w} \cdot \varphi(\vec{x}_i) - y_i &= \left[\sum_{j=1}^n c_j y_j \varphi(\vec{x}_j) \cdot \varphi(\vec{x}_i)\right] - y_i \\ &= \left[\sum_{j=1}^n c_j y_j k(\vec{x}_j, \vec{x}_i)\right] - y_i. \end{aligned}$$

Multi class problems:






Multi class problems: 1-vs-Rest



N models, take best.

Support Vector Machines

Multi class problems: 1-vs-1



$N(N-1)/2$ models – tree execution, take best.

Motivation:

- Modeling a learning problem, we have many parameters to set:
 - Feature extraction algorithm
 - Feature selection and reduction
 - Choice of the learning algorithm
 - Parameters of the learning algorithm
 - ...
- **How to find the best model ?**

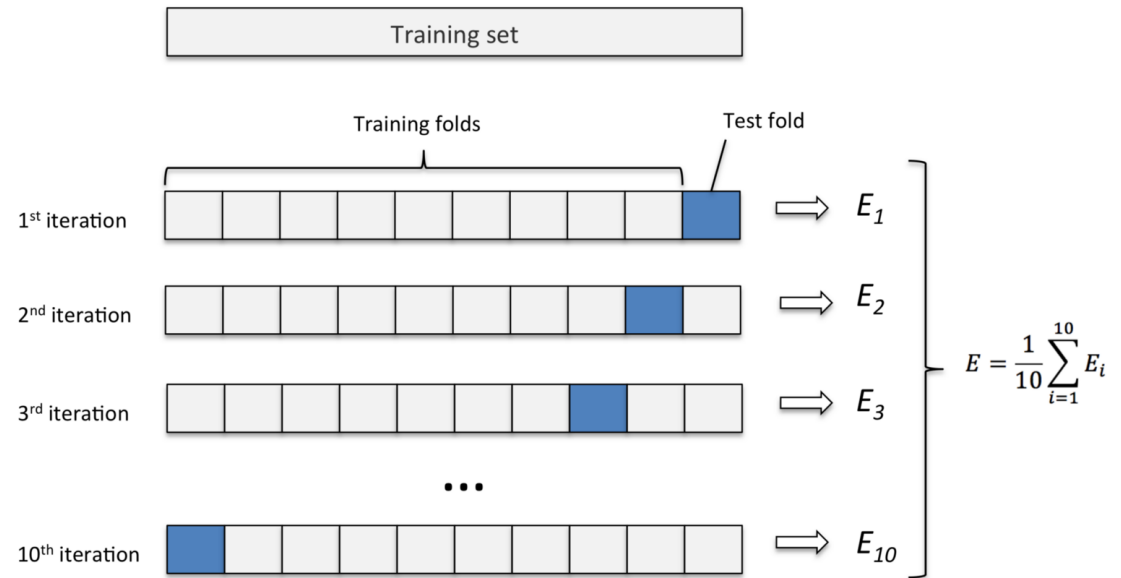
How to find the best model ?

- We already have quality measures to compare models, like
 - Accuracy
 - F-Measure
 - ROC
 - ...
- **But there are two problems:**
 - Do not over fit on the test data (→ can't test too often to stay unbiased)
 - Computational complexity (→ can't try every possible combination)

Tuning a model without test data

Simple approach: **n-fold cross validation**

- Split train data into n parts
- Train on n-1 parts – test on the left out part
- Repeat n-times, leaving out a different part each time
- Average test results



How to find the best model ?

- Grid-search over the parameter space
 - Very expensive
 - How to space the grid ?
- **Random Search**
 - Cheaper than grid-search
 - Quite effective
- Bayesian optimization
 - “optimal” next parameter set for testing

Lab exercises coming up ...